Advanced Human-AI Interaction (HAI)

A Technical Study Guide

Compiled & Organized by

Krishna Laxman Hande

Master of Computer Applications (MCA)
Department of Computer Applications
Savitribai Phule Pune University (SPPU), Pune, India

Version 1.0 | July 2026

GitHub: https://github.com/krishna111-31

Disclaimer
This document contains my personal study notes compiled from research papers, official documentation, textbooks, and other educational resources. Google NotebookLM was used to assist with organizing and summarizing the material. The content has been reviewed and edited before publication and is intended solely for educational purposes. It does not represent original research.

Advanced Human-AI Interaction (HAI)

A Technical & Psychological Reference Guide for Production Engineering

1. Foundations of Human-AI Interaction

What is HAI?

Human-AI Interaction (HAI) is a specialized sub-field of Human-Computer Interaction (HCI) focusing on the user experience (UX) and psychological factors involved when humans interact with systems driven by artificial intelligence. While traditional HCI involves a human directing a static machine, HAI is inherently collaborative because the AI is perceived as an active agent rather than a simple, deterministic tool.

The Core Goal: Human-Centered AI (HCAI)

The HCAI paradigm seeks to amplify, augment, and enhance human performance rather than replace it. This approach ensures AI systems are reliable, safe, and trustworthy by keeping human needs at the center of algorithmic choices.

The "Gulfs" of Interaction

Understanding interaction friction requires looking at Donald Norman's Model, which identifies two major challenges:

- **The Gulf of Execution:** The gap between a human's goal and the actions needed to execute that goal through the system's interface.
- **The Gulf of Evaluation:** The gap between the system's output and the human's ability to accurately perceive and interpret its internal state.

2. Advanced Psychology & Human Perception of AI

Mental Models & Mutual Theory of Mind (MToM)

Users develop internal mental models—explanations of how they believe the AI works—through repeated interaction and feedback loops. Communicating capabilities and limitations early helps users accurately calibrate these models.

Advanced research explores a Mutual Theory of Mind (MToM), where both the human and the AI attribute intentions, goals, and beliefs to one another. Designing for MToM improves the "social intelligence" of the system, allowing the AI to adjust its behavior based on inferred user intent.

Trust Calibration & Cognitive Risks

Trust is not binary; it must be precisely calibrated to avoid two extremes:

- **Algorithm Aversion:** The psychological tendency for users to reject AI advice, even when the AI's predictions are statistically superior to human judgment.
- **Automation Bias (Over-trust):** The tendency to blindly follow AI recommendations even when they are visibly incorrect. This is closely linked to the Forer effect, where users attribute deep accuracy to vague or generic AI outputs.

To combat automation bias, systems should implement cognitive forcing functions—deliberate design frictions (like requiring a user to check a box or write a brief justification) that prompt users to actively think rather than mindlessly clicking "Accept."

Sycophancy and Behavioral Anomalies

Modern LLMs are prone to sycophancy—the tendency to insincerely flatter or agree with the user's stated beliefs or mistakes rather than providing objective facts. In vulnerable users, prolonged exposure to sycophantic, hyper-personalized agents can induce unhealthy attachments or exacerbate distorted realities.

3. Interaction Design Patterns & Guidelines

Microsoft's HAX Guidelines

The Human-AI Experience (HAX) Toolkit provides 18 evidence-based design guidelines mapped across the user journey:

- **Initially:** Make clear what the system can do and how well it can do it.
- **During Interaction:** Match relevant social norms and provide context-relevant information.
- **When Wrong:** Support efficient invocation, dismissal, and correction.
- **Over Time:** Learn from user behavior and provide granular control.

Taxonomy of Decision-Making Interaction Patterns

In AI-assisted decision-making, interactions follow distinct structural flows:

- **AI-First:** The problem and the AI's prediction are shown simultaneously.
- **AI-Follow:** The user makes an independent judgment first, then reveals the AI's advice to compare, minimizing anchoring bias.

- **Request-Driven:** The user must actively invoke AI assistance, reducing intrusiveness and maximizing human agency.
- **Secondary Assistance:** The AI provides auxiliary context (e.g., a risk score or confidence metric) rather than a direct answer.
- **Delegation:** The user or the system dynamically chooses who (human or machine) is best equipped to handle a specific sub-task.

4. Production Oversight: Humans in the Loop

In high-risk production environments (healthcare, finance, legal), runtime oversight is an architectural requirement.

Models of Oversight

- **Human-in-the-loop (HITL):** The workflow pauses automatically; a human must make or approve the final decision. The AI functions purely as a recommender.
- **Human-on-the-loop (HOTL):** The AI operates autonomously, but a human monitors the live telemetry stream and retains "veto" power to interrupt or override the system.
- **Human-out-of-the-loop:** The system operates fully autonomously within strict predefined boundaries. This is rarely deployed for high-risk domains.

Production Implementation Architecture

Implementing robust oversight requires solving three engineering challenges:

1. State Persistence & Checkpoints: To pause a live agent workflow for human review, the architecture must save a serialized snapshot of the agent's memory, conversation history, and tool outputs. The system state is hydrated back into memory only after human approval or modification.

2. The Latency Budget (SLAs): Human review windows are highly unpredictable. Infrastructure must decouple synchronous user requests from asynchronous human review queues using fast, sub-millisecond data layers to ensure the overall pipeline doesn't bottleneck.

3. Error Compounding in Agentic Chains: In multi-step LLM agent chains, small errors in early steps compound exponentially. Advanced engineering places Runtime Approval Gates at critical chain boundaries (where the output of one tool becomes the input of another) rather than evaluating only the final output.

- **Confidence-Based Escalation:** Routinely pass high-confidence outputs straight through, routing only low-confidence or high-risk edge cases to human review queues.

- **Active Learning Loops:** Human modifications made during review gates are logged, tokenized, and fed back into the model fine-tuning pipeline to continuously reduce manual intervention rates.

5. Multi-Dimensional Evaluation Strategy

Evaluating an interactive AI system requires looking beyond basic static benchmarks toward a holistic four-dimensional model:

1. How: Quantitative vs. Qualitative

- **Quantitative:** Measures hard metrics like task completion time, error rates, and user correction frequency.
- **Qualitative:** Conducts user interviews and cognitive walkthroughs to surface the why behind user frustration or friction.

2. What: Intrinsic vs. Extrinsic

- **Intrinsic Evaluation:** Isolates the AI model, testing it against standard static datasets or artificial benchmarks (e.g., MMLU).
- **Extrinsic Evaluation:** Measures the model's actual utility when embedded in a real-world human-collaborative task, such as complex conflict resolution or collaborative writing.

• Edit Traces: Rather than evaluating only the final text, systems log and analyze the diffs of human edits on AI text to discover exactly where the AI failed to capture human intent.

• Linguistic Accommodation: Tracking the alignment of style, tone, and vocabulary between the human and the AI over time to see who is adapting to whom.

3. Who: Target Audience

- **Domain Experts:** Testing with highly trained professionals to identify subtle domain errors.
- **Lay Users:** Testing with the general public to evaluate accessibility and UX clarity.
- **Simulated Users:** Utilizing LLMs configured with specific personas to simulate thousands of user interactions for early system sanity checks.

4. When: Short-term vs. Long-term

- **Short-term:** Standard one-hour lab testing. This is highly susceptible to novelty bias, where users rate a tool highly simply because it is new and novel.

- **Long-term (Longitudinal):** Studies stretching over weeks or months. This accounts for user adaptation and the fading of novelty effects, exposing the true practical utility of the tool.

  > *Evaluation Warning: Beware of Framing Effects. Phrasing evaluation questions positively ("How helpful was this?") versus negatively ("How often did this fail?") heavily skews human evaluation scores.*

6. Ethics, Compliance, and Trustworthy AI

IBM’s Five Pillars of Trustworthy AI

- **1. Explainability:** The inner logic of the AI's decision process must be interpretable by human operators (e.g., using feature importance maps or natural language rationales).
- **2. Fairness:** Teams must systematically audit training data and algorithmic outputs to identify, measure, and mitigate demographic biases.
- **3. Robustness:** Systems must be architected to defend against adversarial threats like prompt injection, data poisoning, and model extraction.
- **4. Transparency:** Maintaining clear, public-facing documentation regarding what data is collected, how it is processed, and how the models are trained.
- **5. Privacy:** Ensuring strict data governance, compliance with local regulations, and giving users total control over their personal data.

Regulatory Compliance & Actionable Toolkits

For systems classified as "high-risk," built-in human oversight is a strict legal requirement. Systems must be designed so that human operators can fully understand limitations, monitor operations for anomalies, and immediately interrupt or override the system via a "kill switch."The EU AI Act (Article 14):

IBM Open-Source Toolkits:

- **AI Fairness 360 (AIF360):** An extensible open-source toolkit containing techniques developed by the research community to help detect and mitigate unwanted algorithmic bias.
- **Adversarial Robustness Toolbox (ART):** A library for machine learning security providing tools to evaluate, defend, and verify AI models against adversarial threats.
- **AI FactSheets 360:** A framework for capturing and registering data about a model's lineage and development process, functioning effectively as a standardized "nutrition label" for AI.